

## SQL Leet Code:

### ) Replace Employee ID with Unique Identifier:

#### Table: Employees

Column Name	Type
id	int
name	varchar

id is the primary key (column with unique values) for this table. Each row of this table contains the id and the name of an employee in a company.

#### Table: EmployeeUNI

Column Name	Type
id	int
unique_id	int

(id, unique\_id) is the primary key (combination of columns with unique values) for this table. Each row of this table contains the id and the corresponding unique\_id of an employee in the company.

Write a solution to show the unique\_id of each user. If a user does not have a unique\_id, display null. Return the result table in any order.

#### Example 1:

##### Input:

Employees table:

id	name
1	Alice
7	Bob
11	Meir
90	Winston
3	Jonathan

EmployeeUNI table:

id	unique_id
3	1

11	2
90	3

### Output:

unique_id	name
null	Alice
null	Bob
2	Meir
3	Winston
1	Jonathan

### Explanation:

- Alice and Bob do not have a `unique_id`, so we display `null`.
- The `unique_id` of Meir is 2.
- The `unique_id` of Winston is 3.
- The `unique_id` of Jonathan is 1.

### Solution:

```
SELECT eu.unique_id,e.name
FROM Employees e
LEFT JOIN EmployeeUNI eu
ON e.id=eu.id;
```

### 2) Top Travellers:

#### Table: Users

Column Name	Type
id	int
name	varchar

`id` is the column with unique values for this table. `name` is the name of the user.

#### Table: Rides

Column Name	Type
id	int
user_id	int
distance	int

`id` is the column with unique values for this table. `user_id` is the id of the user who traveled the distance "distance".

Write a solution to report the distance traveled by each user. Return the result table ordered by `travelled_distance` in descending order. If two or more users traveled the same distance, order them by their `name` in ascending order. The result format is in the following example.

### Example 1:

#### Input:

Users table:

id	name
1	Alice
2	Bob
3	Alex
4	Donald
7	Lee
13	Jonathan
19	Elvis

Rides table:

id	user_id	distance
1	1	120
2	2	317
3	3	222
4	7	100
5	13	312
6	19	50
7	7	120
8	19	400
9	7	230

#### Output:

name	travelled_distance
Elvis	450
Lee	450
Bob	317
Jonathan	312
Alex	222
Alice	120
Donald	0

**Explanation:**

Elvis and Lee traveled 450 miles. Elvis is the top traveler as his name is alphabetically smaller than Lee. Bob, Jonathan, Alex, and Alice have only one ride and we just order them by the total distances of the ride. Donald did not have any rides, the distance traveled by him is 0.

**Solution:**

```
SELECT u.name,  
NVL(SUM(r.distance), 0) AS travelled_distance  
FROM Users u  
LEFT JOIN Rides r ON u.id = r.user_id  
GROUP BY u.name  
ORDER BY travelled_distance DESC, u.name ASC;
```

**3) Group Sold Products By The Date:**

**Table:** Activities

Column Name	Type
sell_date	date
product	varchar

There is no primary key (column with unique values) for this table. It may contain duplicates. Each row of this table contains the product name and the date it was sold in a market.

Write a solution to find for each date the number of different products sold and their names. The sold products names for each date should be sorted lexicographically. Return the result table ordered by `sell_date`. The result format is in the following example.

**Example 1:****Input:**

Activities table:

sell_date	product
2020-05-30	Headphone
2020-06-01	Pencil
2020-06-02	Mask
2020-05-30	Basketball
2020-06-01	Bible

2020-06-02	Mask
2020-05-30	T-Shirt

### Output:

sell_date	num_sold	products
2020-05-30	3	Basketball,Headphone,T-Shirt
2020-06-01	2	Bible,Pencil
2020-06-02	1	Mask

### Explanation:

- For 2020-05-30, sold items were <sup>1</sup> (Basketball, Headphone, T-Shirt). We sort them lexicographically and separate them by a comma.
- For 2020-06-01, sold items were (Bible, Pencil). We sort them lexicographically and separate them by a comma.
- For 2020-06-02, the sold item is (Mask). We just return it.

### Solution:

SELECT

sell\_date,

COUNT(DISTINCT product) AS num\_sold,

LISTAGG(DISTINCT product, ',') WITHIN GROUP (ORDER BY product)  
AS products

FROM Activities

GROUP BY sell\_date

ORDER BYsell\_date;

### 4)

**Table:** Visits

Column Name	Type
visit_id	int
customer_id	int

visit\_id is the column with unique values for this table. This table contains information about the customers who visited the mall.

**Table: Transactions**

Column Name	Type
transaction_id	int
visit_id	int
amount	int

transaction\_id is the column with unique values for this table. This table contains information about the transactions made during the visit\_id.

Write a solution to find the IDs of the users who visited without making any transactions and the number of times they made these types of visits. Return the result table sorted in any order. The result format is in the following example.

**Example 1:****Input:**

Visits table:

visit_id	customer_id
1	23
2	9
4	30
5	54
6	96
7	54
8	54

Transactions table:

transaction_id	visit_id	amount
2	5	310
3	5	300
9	5	200
12	1	910
13	2	970

**Output:**

customer_id	count_no_trans
54	2
30	1
96	1

**Explanation:**

- Customer with <sup>1</sup>customer\_id = 23 visited the mall once and made one transaction during the visit with visit\_id = 12.
- Customer with customer\_id = 9 visited the mall once and made one transaction during the visit with visit\_id = 13.
- Customer with customer\_id = 30 visited the mall once and did not make any transactions.
- Customer with customer\_id = 54 visited the mall three times. During 2 visits they did not make any transactions, and during one visit they made 3 transactions.
- Customer with customer\_id = 96 visited the mall once and did not make any transactions.

**Solution:**

```
SELECT v.customer_id, COUNT (v.visit_id) AS  
count_no_trans  
FROM Visits v  
LEFT JOIN Transactions t  
ON v.visit_id = t.visit_id  
WHERE t.transaction_id IS NULL  
GROUP BY v.customer_id;
```

5)

**Table:** Users

Column Name	Type
account	int
name	varchar

account is the primary key (column with unique values) for this table. Each row of this table contains the account number of each user in the bank. There will be no two users having the same name in the table.

**Table:** Transactions

Column Name	Type
-------------	------

trans_id	int
account	int
amount	int
transacted_on	date

trans\_id is the primary key (column with unique values) for this table. Each row of this table contains all changes made to all accounts. amount is positive if the user received money and negative if they transferred money. All accounts start with a balance of 0.

Write a solution to report the name and balance of users with a balance higher than 10000. The balance of an account is equal to the sum of the amounts of all transactions involving that account. Return the result table in any order. The result format is in the following example.

### Example 1:

#### Input:

Users table:

account	name
900001	Alice
900002	Bob
900003	Charlie

Transactions table:

trans_id	account	amount	transacted_on
1	900001	7000	2020-08-01
2	900001	7000	2020-09-01
3	900001	-3000	2020-09-02
4	900002	1000	2020-09-12
5	900003	6000	2020-08-07
6	900003	6000	2020-09-07
7	900003	-4000	2020-09-11

#### Output:

name	balance
Alice	11000

#### Explanation:

- Alice's balance is  $(7000 + 7000 - 3000) = 11000$ .



- Bob's balance is 1000.
- Charlie's balance is  $(6000 + 6000 - 4000) = 8000$ .

Solution:

```
SELECT u.name, SUM amount) AS balance
FROM Transactions t
INNER JOIN users u
ON t.account = u.account
GROUP BY u.name
HAVING SUM amount > 10000;
```

6)

**Table:** Tweets

Column Name	Type
tweet_id	int
content	varchar

tweet\_id is the primary key (column with unique values) for this table.<sup>1</sup> content consists of alphanumeric characters, '!', or ' ' and no other special characters. This table contains all the tweets in a social media app.

Write a solution to find the IDs of the invalid tweets. The tweet is invalid if the number of characters used in the content of the tweet is strictly greater than 15. Return the result table in any order. The result format is in the following example.

**Example 1:**

**Input:**

Tweets table:

tweet_id	content
1	Let us Code
2	More than fifteen chars are here!

**Output:**

```

+-----+
| tweet_id |
+-----+
| 2        |
+-----+

```

**Explanation:**

- Tweet 1 has length = 11. It is a valid tweet.
- Tweet 2 has length = 33. It is an invalid tweet.

**Solution:**

```

SELECT tweet_id
FROM Tweets
WHERE LENGTH(content) > 15;

```

7)

**Table: Employees**

```

+-----+-----+
| Column Name | Type |
+-----+-----+
| emp_id      | int  |
| event_day   | date |
| in_time     | int  |
| out_time    | int  |
+-----+-----+

```

(emp\_id, event\_day, in\_time) is the primary key<sup>1</sup> (combinations of columns with unique values) of this table. The table shows the employees' entries and exits in an office. event\_day is the day at which this event happened, in\_time is the minute at which the employee entered the office, and out\_time is the minute at which they left the office. in\_time and out\_time are between 1 and 1440. It is guaranteed that no two events on the same day intersect in time, and in\_time < out\_time.

Write a solution to calculate the total time (in minutes) spent by each employee on each day at the office. Note that within one day, an employee can enter and leave more than once. The time spent in the office for a single entry is out\_time - in\_time. Return the result table in any order. The result format is in the following example.

**Example 1:****Input:**

Employees table:

emp_id	event_day	in_time	out_time
1	2020-11-28	4	32
1	2020-11-28	55	200
1	2020-12-03	1	42
2	2020-11-28	3	33
2	2020-12-09	47	74

### Output:

day	emp_id	total_time
2020-11-28	1	173
2020-11-28	2	30
2020-12-03	1	41
2020-12-09	2	27

### Explanation:

- Employee 1 has three events: two on day 2020-11-28 with a total of  $(32 - 4) + (200 - 55) = 173$ , and one on day 2020-12-03 with a total of  $(42 - 1) = 41$ .
- Employee 2 has two events: one on day 2020-11-28 with a total of  $(33 - 3) = 30$ , and one on day 2020-12-09 with a total of  $(74 - 47) = 27$ .

### Solution:

```
SELECT event_day AS day, emp_id, SUM out_time_in_time AS total_time
FROM Employees
GROUP BY event_day, emp_id;
```

8)

### Table: Products

Column Name	Type
product_id	int
low_fats	enum
recyclable	enum

product\_id is the primary key (column with unique<sup>1</sup> values) for this table. low\_fats is an ENUM (category) of type ('Y', 'N') where 'Y' means this product is low fat and 'N' means it is

not. `recyclable` is an ENUM (category) of types ('Y', 'N') where 'Y' means this product is recyclable and 'N' means it is not.

Write a solution to find the ids of products that are both low fat and recyclable. Return the result table in any order. The result format is in the following example.

### Example 1:

#### Input:

Products table:

product_id	low_fats	recyclable
0	Y	N
1	Y	Y
2	N	Y
3	Y	Y
4	N	N

#### Output:

product_id
1
3

#### Explanation:

Only products 1 and 3 are both low fat and recyclable.

Solution:

```
SELECT product_id
FROM Products
WHERE low_fats = 'Y' AND recyclable = 'Y';
```

10)

**Table:** Employees

Column Name	Type
employee_id	int
name	varchar
salary	int

employee\_id is the primary key (column with unique values) for this table. Each row of this table indicates the employee ID, employee name, and salary.

Write a solution to calculate the bonus of each employee. The bonus of an employee is 100% of their salary if the ID of the employee is an odd number and the employee's name does not start with the character 'M'. The bonus of an employee is 0 otherwise. Return the result table ordered by employee\_id. The result format is in the following example.

### Example 1:

#### Input:

Employees table:

employee_id	name	salary
2	Meir	3000
3	Michael	3800
7	Addilyn	7400
8	Juan	6100
9	Kannon	7700

#### Output:

employee_id	bonus
2	0
3	0
7	7400
8	0
9	7700

#### Explanation:

- The employees with IDs 2 and 8 get 0 bonus because they have an even employee\_id.
- The employee with ID 3<sup>1</sup> gets 0 bonus because their name starts with 'M'.
- The rest of the employees get a 100% bonus.

#### Solution:

```

SELECT employee_id, CASE WHEN employee_id % 2 <> 0 AND name NOT
LIKE 'M%' THEN salary
ELSE 0 END AS bonus
FROM Employees
ORDER BY employee_id;

```

11)

**Table:** Logins

Column Name	Type
user_id	int
time_stamp	datetime

(user\_id, time\_stamp) is the primary key (combination of columns with unique values) for this table. Each row contains information about the login time for the user with ID user\_id.

Write a solution to report the latest login for all users in the year 2020. Do not include the users who did not login in 2020. Return the result table in any order. The result format is in the following example.

**Example 1:**

**Input:**

Logins table:

user_id	time_stamp
6	2020-06-30 15:06:07
6	2021-04-21 14:06:06
6	2019-03-07 00:18:15
8	2020-02-01 05:10:53
8	2020-12-30 00:46:50
2	2020-01-16 02:49:50
2	2019-08-25 07:59:08
14	2019-07-14 09:00:00
14	2021-01-06 11:59:59

**Output:**

user_id	last_stamp
6	2020-06-30 15:06:07

8	2020-12-30 00:46:50
2	2020-01-16 02:49:50

### Explanation:

- User 6 logged into their account 3 times but only once in 2020, so we include this login in the result table.
- User 8 logged into their account 2 times in 2020, once in February and once in December. We include only the latest one (December) in the result table.
- User 2 logged into their account 2 times but only once in 2020, so we include this login in the result table.
- User 14 did not login in 2020, so we do not include them in the result table.

### Solution:

```
SELECT DISTINCT user_id, FIRST_VALUE time_stamp OVER PARTITION
BY user_id ORDER BY time_stamp DESC AS last_stamp
FROM Logins
WHERE YEAR time_stamp = '2020';
```

### 12)

#### Table: Employees

Column Name	Type
employee_id	int
name	varchar

employee\_id is the column with unique values for this table. Each row of this table indicates the name of the employee whose ID is employee\_id.

#### Table: Salaries

Column Name	Type
employee_id	int
salary	int

employee\_id is the column with unique values for this table. Each row of this table indicates the salary of the employee whose ID is employee\_id.

Write a solution to report the IDs of all the employees with missing information. The information of an employee is missing if:

- The employee's `name` is missing, or
- The employee's `salary` is missing.

Return the result table ordered by `employee_id` in ascending order. The result format is in the following example.

### Example 1:

#### Input:

Employees table:

employee_id	name
2	Crew
4	Haven
5	Kristian

Salaries table:

employee_id	salary
5	76071
1	22517
4	63539

#### Output:

employee_id
1
2

#### Explanation:

Employees 1, 2, 4, and 5 are working at this company.

- The name of employee 1 is missing.
- The salary of employee 2<sup>1</sup> is missing.

Solution:

```
SELECT e.employee_id
```

```
FROM Employees e
```



```

LEFT JOIN Salaries s
ON e.employee_id = s.employee_id
WHERE s.salary IS NULL
UNION
SELECT s.employee_idFROM Salaries s
LEFT JOIN Employees e
ON s.employee_id = e.employee_id
WHERE e.name IS NULL
ORDER BY employee_id;

```

13)

**Table: Stocks**

Column Name	Type
stock_name	varchar
operation	enum
operation_day	int
price	int

(stock\_name, operation\_day) is the primary key (combination of columns with unique values) for this table. The operation column is an ENUM (category) of type ('Sell', 'Buy'). Each row of this table indicates that the stock which has stock\_name had an operation on the day operation\_day with the price. It is guaranteed that each 'Sell' operation for a stock has a corresponding 'Buy' operation in a previous day. It is also guaranteed that each 'Buy' operation for a stock has a corresponding 'Sell' operation in an upcoming day.

Write a solution to report the Capital gain/loss for each stock. The Capital gain/loss of a stock is the total gain or loss after buying and selling the stock one or many times. Return the result table in any order. The result format is in the following example.

**Example 1:**

**Input:**

Stocks table:

stock_name	operation	operation_day	price
------------	-----------	---------------	-------

Leetcode	Buy	1	1000	
Corona Masks	Buy	2	10	
Leetcode	Sell	5	9000	
Handbags	Buy	17	30000	
Corona Masks	Sell	3	1010	
Corona Masks	Buy	4	1000	
Corona Masks	Sell	5	500	
Corona Masks	Buy	6	1000	
Handbags	Sell	29	7000	
Corona Masks	Sell	10	10000	
+-----+-----+-----+-----+				

### Output:

+-----+-----+	
stock_name	capital_gain_loss
+-----+-----+	
Corona Masks	9500
Leetcode	8000
Handbags	-23000
+-----+-----+	

### Explanation:

- Leetcode stock was bought at day 1 for 1000\$ and was sold at day 5 for 9000\$.  
Capital gain = 9000 - 1000 = 8000\$.
- Handbags stock was bought at day 17 for 30000\$ and was sold at day 29 for 7000\$.  
Capital loss = 7000 - 30000 = -23000\$.
- Corona Masks stock was bought at day 1 for 10\$ and was sold at day 3 for 1010\$. It was bought again at day 4 for 1000\$ and was sold at day 5 for 500\$. At last, it was bought at day 6 for 1000\$ and was sold at day 10 for 10000\$. Capital gain/loss is the sum of capital gains/losses for each ('Buy' --> 'Sell') operation = (1010 - 10) + (500 - 1000) + (10000 - 1000) = 1000 - 500 + 9000 = 9500\$.

### Solution:

```
SELECT stock_name, SUM CASE WHEN operation = 'Buy' THEN price*-1
ELSE price END AS capital_gain_loss
```

```
FROM Stocks
```

```
GROUP BY stock_name;
```

### 14)

#### Table: Products

+-----+-----+	
Column Name	Type
+-----+-----+	
product_id	int
store1	int

store2	int	
store3	int	
+-----+-----+		

`product_id` is the primary key (column with unique values) for this table. Each row in this table indicates the product's price in 3 different stores: `store1`, `store2`, and `store3`. If the product is not available in a store, the price will be null in that store's column.

Write a solution to rearrange the `Products` table so that each row has (`product_id`, `store`, `price`). If a product is not available in a store, do not include a row with that `product_id` and `store` combination in the result table. Return the result table in any order. The result format is in the following example.

### Example 1:

#### Input:

`Products` table:

+-----+-----+-----+-----+
product_id   store1   store2   store3
+-----+-----+-----+-----+
0            95       100      105
1            70       null     80
+-----+-----+-----+-----+

#### Output:

+-----+-----+-----+
product_id   store    price
+-----+-----+-----+
0            store1   95
0            store2   100
0            store3   105
1            store1   70
1            store3   80
+-----+-----+-----+

#### Explanation:

- Product 0 is available in all three stores with prices 95, 100, and 105 respectively.
- Product 1 is available in store1 with price 70 and store3 with price 80. The product is not available in store2

#### Solution:

```
SELECT product_id, 'store1' AS store, store1 AS price FROM Products
WHERE store1 IS NOT NULL
UNION
```

```

SELECT product_id, 'store2' AS store, store2 AS price FROM Products
WHERE store2 IS NOT NULL

UNION

SELECT product_id, 'store3' AS store, store3 AS price FROM Products
WHERE store3 IS NOT NULL;

```

15)

**Table:** NPV

Column Name	Type
id	int
year	int
npv	int

(id, year) is the primary key of this table. This table contains the ID and the year of the investment and its net present value.

**Table:** Queries

Column Name	Type
id	int
year	int

(id, year) is the primary key of this table. This table contains the ID and the year of the investment.

Write an SQL query to find the npv of each query. Return the result table in any order. The result format is in the following example.

**Example 1:**

**Input:**

NPV table:

id	year	npv
1	2018	100
1	2020	100
2	2020	200

3	2019	200
---	------	-----

Queries table:

id	year
1	2018
2	2020
3	2019
1	2021

**Output:**

id	year	npv
1	2018	100
2	2020	200
3	2019	200
1	2021	0

**Explanation:**

The npv of the query (1, 2021) is not in the NPV table, so its npv is 0.

**Solution:**

SELECT q.id, q.year, NVL(n.npv, 0) AS npv

FROM Queries q

LEFT JOIN

NPV n ON q.id = n.id AND q.year = n.year;

## 16) Average Selling Price

**Table:** Products

Column Name	Type
product_id	int
product_name	varchar
unit_price	int

product\_id is the primary key of this table. Each row of this table indicates the name and the price of each product.

**Table:** Orders

Column Name	Type
product_id	int
quantity	int
order_date	date

(product\_id, order\_date) is the primary key of this table. Each row of this table indicates the quantity of product\_id was ordered on order\_date.

Write an SQL query to find the average selling price for each product. The average\_price should be rounded to 2 decimal places. Return the result table in any order.

The result format is in the following example.

### Example 1:

#### Input:

Products table:

product_id	product_name	unit_price
1	LC Phone	300
2	LC T-Shirt	10
3	LC Book	20
4	LC Mug	25

Orders table:

product_id	quantity	order_date
1	2	2020-02-05
1	1	2020-02-10
2	3	2020-02-10
3	1	2020-02-15
1	5	2020-02-25
4	6	2020-02-25

#### Output:

product_id	average_price
1	300.00
2	10.00
3	20.00
4	25.00

#### Explanation:

- Average selling price for product 1 =  $(300 * 2 + 300 * 1 + 300 * 5) / 8 = 300.00$
- Average selling price for product 2 =  $(10 * 3) / 3 = 10.00$
- Average selling price for product 3 =  $(20 * 1) / 1 = 20.00$
- Average selling price for product 4 =  $(25 * 6) / 6 = 25.00$

Solution:

```
SELECT p.product_id,
ROUND(SUM(p.unit_price * o.quantity) / SUM(o.quantity), 2) AS
average_price
FROM Products p
JOIN
Orders o ON p.product_id = o.product_id
GROUP BY p.product_id;
```

## 17) Customer Visits and Purchases

**Table:** Orders

Column Name	Type
order_id	int
order_date	date
customer_id	int
product_id	int

order\_id is the primary key for this table. This table contains<sup>1</sup> the ID of an order, the date of the order, the ID of the customer who ordered it, and the ID of the product which they ordered.

[1. github.com](https://github.com)

[github.com](https://github.com)

Write an SQL query to find the number of times each customer visited the mall and the number of products each customer bought. Return the result table in any order.

The result format is in the following example.

**Example 1:**

**Input:**

Orders table:

order_id	order_date	customer_id	product_id
1	2019-01-01	1	1
2	2019-01-01	2	2
3	2019-01-01	1	3
4	2019-01-01	1	4
5	2019-01-02	2	5
6	2019-01-02	3	6

**Output:**

customer_id	count_visits	count_products
1	2	3
2	2	2
3	1	1

**Explanation:**

- Customer 1 visited the mall twice and bought 3 products.
- Customer 2 visited the mall twice and bought 2 products.
- Customer 3 visited the mall once and bought 1 product.

**Solution:**

```
SELECT customer_id, COUNT(DISTINCT order_date) AS count_visits,  
       COUNT(product_id) AS count_products  
FROM Orders  
GROUP BY customer_id  
ORDER BY customer_id;
```

### 18) Most Recent Orders for Each Product

**Table:** Products

Column Name	Type
product_id	int
product_name	varchar



product\_id is the primary key of this table. Each row of this table indicates the name of each product.

**Table: Orders**

Column Name	Type
product_id	int
order_date	date
unit	int

(product\_id, order\_date) is the primary key of this table. Each row of this table indicates the quantity of each product was ordered on each date.

Write an SQL query to find the most recent order(s) of each product. Return the result table ordered by product\_name in ascending order and order\_date in ascending order.

The result format is in the following example.

**Example 1:**

**Input:**

Products table:

product_id	product_name
1	Avocado
2	Blueberry
3	Almond

Orders table:

product_id	order_date	unit
1	2020-01-02	800
1	2020-01-15	900
2	2020-02-10	600
2	2020-03-01	700
3	2020-01-07	750
3	2020-01-30	800
3	2020-03-15	1000

**Output:**

product_name	order_date	unit
Almond	2020-03-15	1000

Avocado	2020-01-15	900	
Blueberry	2020-03-01	700	
+-----+	+-----+	+-----+	+-----+

Solution:

```
SELECT w.name AS warehouse_name,
SUM w.units*p.Width*p.Length*p.Height AS volume
FROM Warehouse w
INNER JOIN Products p
ON w.product_id = p.product_id
GROUP BY w.name;
```

OK, here's the formatted and more readable version of the LeetCode problem description:

### 19) Machines That Process a Task

**Table:** Activity

+-----+	+-----+	+-----+
Column Name	Type	
+-----+	+-----+	+-----+
machine_id	int	
process_id	int	
activity_type	enum	
timestamp	float	
+-----+	+-----+	+-----+

The table shows the user activities for a factory production machine.

- machine\_id is the primary key of this table.
- process\_id is the ID of the process run on the machine.
- activity\_type is an ENUM of type ('start', 'end').
- timestamp is a float representing the current time in seconds.
- 'start' means the machine starts the process at the given timestamp, and 'end' means the machine ends the process at the given timestamp.
- The 'start' timestamp is always less than the 'end' timestamp for every (machine\_id, process\_id) combination.

There is a factory production machine that might run various number of processes. Each process is assigned a unique process\_id, and each process may be executed many times.

Write an SQL query to find the average time each machine takes to complete a process. The time to complete a process is the 'end' timestamp minus the 'start' timestamp. The average

time is calculated by the total time to complete every process on the machine divided by the number of processes.

The result should be in ascending order by `machine_id`. Round the average to 3 decimal places.

The result format is in the following example.

### Example 1:

#### Input:

Activity table:

machine_id	process_id	activity_type	timestamp
0	0	start	0.712
0	0	end	1.520
0	1	start	3.140
0	1	end	4.120
1	0	start	0.550
1	0	end	1.550
1	1	start	0.430
1	1	end	1.420
2	0	start	4.100
2	0	end	4.520
2	1	start	2.500
2	1	end	5.000

#### Output:

machine_id	processing_time
0	0.894
1	0.995
2	2.710

#### Explanation:

- Machine 0:
  - process 0:  $1.520 - 0.712 = 0.808$
  - process 1:  $4.120 - 3.140 = 0.980$
  - average:  $(0.808 + 0.980) / 2 = 0.894$
- Machine 1:
  - process 0:  $1.550 - 0.550 = 1.000$
  - process 1:  $1.420 - 0.430 = 0.990$
  - average:  $(1.000 + 0.990) / 2 = 0.995$
- Machine 2:
  - process 0:  $4.520 - 4.100 = 0.420$
  - process 1:  $5.000 - 2.500 = 2.500$

- average:  $(0.420 + 2.500) / 2 = 2.710$

Solution:

```
SELECT problem_id
FROM Problems
WHERE likes / likes+dislikes < 0.6
ORDER BY problem_id;
```

## 20) Number of Products Ordered in the Period

**Table:** Products

Column Name	Type
product_id	int
product_name	varchar

product\_id is the primary key of this table. Each row of this table indicates the name of each product.

**Table:** Orders

Column Name	Type
product_id	int
order_date	date
unit	int

(product\_id, order\_date) is the primary key of this table. Each row of this table indicates the quantity of each product was ordered on each date.

Write an SQL query to find the number of products that were ordered in the year 2020.  
Return the result table in any order.

The result format is in the following example.

**Example 1:**

**Input:**

Products table:

product_id	product_name
1	Leetcode Mobile
2	Google Pixel
3	Samsung Galaxy

Orders table:

product_id	order_date	unit
1	2020-02-05	60
2	2020-03-01	70
3	2020-04-18	80
1	2020-02-09	100
2	2020-06-30	90
3	2020-12-25	200
1	2019-01-01	50
2	2019-02-02	60
3	2019-03-03	70

**Output:**

products_count
3

**Explanation:**

All products were ordered in the year 2020.

**Solution:**

```
SELECT COUNT(DISTINCT product_id) AS products_count
FROM Orders
WHERE EXTRACT(YEAR FROM order_date) = 2020;
```

## 21) Daily Active Users II

**Table:** Activities

Column Name	Type
activity_date	date

user_id	int	
activity_type	enum	
+-----+		

There is no primary key for this table; it may have duplicate rows. The `activity_type` column is an ENUM of type ('open\_session', 'end\_session', 'scroll\_down', 'send\_message'). The table shows the user activities for a social media website.

Write an SQL query to find the daily active user count for a period of 30 days ending 2019-07-27 inclusively. A user was active on a day if they made at least one activity on that day.

Return the result table in any order.

The query result format is in the following example.

### Example 1:

#### Input:

Activities table:

activity_date	user_id	activity_type	
+-----+			
2019-07-27	1	open_session	
2019-07-27	2	open_session	
2019-07-25	1	end_session	
2019-07-25	2	scroll_down	
2019-07-25	1	send_message	
2019-07-25	3	open_session	
2019-07-25	3	end_session	
2019-07-24	1	open_session	
2019-07-24	1	end_session	
2019-07-24	2	scroll_down	
2019-07-23	1	send_message	
2019-07-23	2	open_session	
2019-07-23	3	scroll_down	
2019-07-23	3	send_message	
2019-06-23	4	open_session	
+-----+			

#### Output:

activity_date	user_count	
+-----+		
2019-07-27	2	
2019-07-25	3	
2019-07-24	2	
2019-07-23	3	
+-----+		

#### Explanation:

Note that we only care about dates with activity in the range between 2019-06-28 and 2019-07-27 inclusive.

- On 2019-07-27, users 1 and 2 were active.
- On 2019-07-25, users 1, 2, and 3 were active.
- On 2019-07-24, users 1 and 2 were active.
- On 2019-07-23, users 1, 2, and 3 were active.

Solution:

```
SELECT activity_date, COUNT(DISTINCT user_id) AS user_count
FROM Activities
WHERE activity_date BETWEEN ADD_MONTHS(Date '2019-07-27', -1) +
1 AND Date '2019-07-27'
GROUP BY activity_date;
```

## 22) Customer Who Placed the Largest Number of Orders

**Table:** Orders

Column Name	Type
order_id	int
customer_name	varchar
customer_id	int

order\_id is the primary key for this table. customer\_name is the name of the customer who placed the order. customer\_id is the ID of the customer who placed the order.

**Table:** Customers

Column Name	Type
customer_id	int
customer_name	varchar

customer\_id is the primary key for this table. customer\_name is the name of the customer.

Write an SQL query to find the name of the customer who has placed the most orders.

Return the result table in any order.

The result format is in the following example.

### Example 1:

#### Input:

Orders table:

order_id	customer_name	customer_id
1	Diana	1
2	Diana	1
3	Nicholas	2
4	Diana	1
5	Harrison	3
6	Alice	4

Customers table:

customer_id	customer_name
1	Diana
2	Nicholas
3	Harrison
4	Alice

#### Output:

customer_name
Diana

#### Explanation:

Diana has placed 3 orders, which is the most orders any customer has placed.

Solution:

```
SELECT customer_name
FROM ( SELECT customer_name,
              COUNT(*) AS order_count
        FROM Orders
        GROUP BY customer_name
        ORDER BY COUNT(*) DESC)
```



WHERE ROWNUM = 1;

### 23) Activity With the Most Users

**Table:** Friends

Column Name	Type
id	int
name	varchar
activity	varchar

id is the primary key for this table. name is the name of the user. activity is the name of the activity the user likes to do.

**Table:** Activities

Column Name	Type
activity	varchar

activity is the primary key for this table. activity is the name of some activity.

Write an SQL query to find the activity name with the maximum number of distinct users that like it.

Return the result table in any order.

The result format is in the following example.

#### Example 1:

##### Input:

Friends table:

id	name	activity
1	Jonathan	Hiking
2	Amy	Hiking
3	Drew	Swimming
4	Alice	Swimming
5	Daniel	Hiking
6	Helen	Running
7	Jonathan	Running
8	Lisa	Running

Activities table:

```

+-----+
| activity |
+-----+
| Hiking   |
| Swimming |
| Running  |
+-----+

```

### Output:

```

+-----+
| activity |
+-----+
| Hiking   |
+-----+

```

### Explanation:

There are 3 distinct users who like "Hiking", 2 distinct users who like "Swimming", and 3 distinct users who like "Running".

Since "Hiking" and "Running" have the same maximum number of distinct users, we return "Hiking" as it has the smallest lexicographical order.

### Solution:

```

SELECT activity
FROM (
    SELECT activity,
           COUNT(DISTINCT id) AS user_count
    FROM Friends
    GROUP BY activity
    ORDER BY COUNT(DISTINCT id) DESC, activity AS)
WHERE ROWNUM = 1;

```

## 24) Department Total Revenue

**Table:** Departments

```

+-----+-----+
| Column Name | Type |
+-----+-----+
| id          | int  |

```

```

| name          | varchar |
| revenue       | int     |
| month         | varchar |
+-----+-----+

```

(id, month) is the primary key of this table. The month column has values in the format "YYYY-MM".

Write an SQL query to find the total revenue for each department in each year.

Return the result table in any order.

The result format is in the following example.

### Example 1:

#### Input:

Departments table:

```

+-----+-----+-----+-----+
| id    | name  | revenue | month  |
+-----+-----+-----+-----+
| 1     | A     | 7000    | 2018-01 |
| 2     | A     | 6500    | 2018-02 |
| 3     | A     | 8000    | 2018-03 |
| 4     | B     | 5000    | 2018-01 |
| 5     | B     | 6000    | 2018-02 |
| 6     | C     | 4000    | 2018-01 |
| 7     | A     | 9000    | 2019-01 |
| 8     | A     | 7000    | 2019-02 |
| 9     | B     | 6000    | 2019-03 |
+-----+-----+-----+-----+

```

#### Output:

```

+-----+-----+-----+
| dept_name | year  | total_revenue |
+-----+-----+-----+
| A         | 2018  | 21500          |
| A         | 2019  | 16000          |
| B         | 2018  | 11000          |
| B         | 2019  | 6000           |
| C         | 2018  | 4000           |
+-----+-----+-----+

```

#### Explanation:

The total revenue for department 'A' in year 2018 is  $7000 + 6500 + 8000 = 21500$ . The total revenue for department 'A' in year 2019 is  $9000 + 7000 = 16000$ . The total revenue for department 'B' in year 2018 is  $5000 + 6000 = 11000$ . The total revenue for department 'B' in year 2019 is 6000. The total revenue for department 'C' in year 2018 is 4000.

Solution:

```
SELECT name AS dept_name,  
       EXTRACT(YEAR FROM TO_DATE(month, 'YYYY-MM')) AS year,  
       SUM(revenue) AS total_revenue  
FROM Departments  
GROUP BY name, EXTRACT(YEAR FROM TO_DATE(month, 'YYYY-MM'))  
ORDER BY name, year;
```

## 25) Consecutive Numbers

**Table:** Logs

Column Name	Type
log_id	int

log\_id is the primary key for this table. Each row of this table contains the ID of a log from a certain database. Since some of the log\_id's were removed from the database, you have missing log\_id's. Consecutive log\_id's are those log\_id's with the difference of 1 between them.

Write an SQL query to find the start and end number of continuous ranges in the table Logs.

Return the result table ordered by start\_id.

The result format is in the following example.

### Example 1:

**Input:**

Logs table:

log_id
1
2
3
7
8
10

+-----+

### Output:

start_id	end_id
1	3
7	8
10	10

### Explanation:

The <sup>1</sup> ranges are:

[1. github.com](#)

[github.com](#)

- 1, 2, 3
- 7, 8
- 10, 10

Solution:

```
SELECT MIN(log_id) AS start_id, MAX(log_id) AS end_id
FROM (
  SELECT log_id, log_id - ROWNUM AS grp
  FROM (
    SELECT log_id
    FROM Logs
    ORDER BY log_id
  )
)
GROUP BY grp
ORDER BY start_id;
```

## 26) Customers With Positive Revenue in the Year 2021

**Table:** Customer

Column Name	Type
customer_id	int
year	int
revenue	int

(customer\_id, year) is the primary key for this table. There are no nulls in any column. The revenue column gives the revenue of the customer with customer\_id in the year year.

Write an SQL query to report the customer\_id from the Customer table that had positive revenue in the year 2021.

### Example 1:

#### Input:

Customer table:

customer_id	year	revenue
1	2021	50
1	2022	100
2	2021	0
3	2021	-50
3	2022	100
4	2021	200
4	2022	-200

#### Output:

customer_id
1
4

#### Explanation:

Customer 1 and 4 had positive revenue in 2021.

#### Solution:

```
SELECT customer_id
```

FROM Customer

WHERE year = 2021 AND revenue > 0;